

Introduction to Relational Databases

- Bachelor Computer Science, Lille 1 University
- Nov 14th, 2012 (lecture 10/12)
- Today's lecturer: C. Kuttler
- Topic: Introduction to SQL
 - Subqueries:
 - Comparison of operators
 - Aggregate functions for generalized AND/OR in having clause
 - Variable visibility
 - Other definitions of data in SQL
 - Views
 - Generic integrity constraints
 - Access control

116

Having clause: generalized AND

- EVERY(test): true, if test is true for all lines of group
- Supplier whose articles all cost at least 10 euro

```
select c.fid
from catalogue c
where
not exists(select *
            from catalogue c2
            where c2.fid=c.fid and c.prix<=10)
```

118

Equivalence of expressive power

- IN, =ANY, EXISTS have the same expressive power, and can also be expressed through a join
- NOT IN, <>ALL, NOT EXISTS have the same expressive power, and can be expressed by a difference
- *comp* SOME, if there are no duplicates, can be rewritten as theta-joins (not as equi-joins)
- *comp* ALL can be rewritten by queries combining grouping and extraction of a minimum and maximum

117

Having clause: generalized AND

- EVERY(test): true, if test is true for all lines of group
 - True for line 1 AND true for line 2 AND ... AND true for line n
 - Other syntax: BOOL_AND(test)
- Supplier whose articles all cost at least 10 euro

```
select fid
from catalogue
group by fid
having every(prix >10)
```

119

Having clause: generalized OR

- `bool_or(test)`: true, if test is true for at least one line of group
- Supplier offering both green and red articles

```
select c.fid
from catalogue c
where
    exists (select * from articles a where
            a.aid=c.aid and a.acoul='green')
and
exists (select * from articles a where
        a.aid=c.aid and a.acoul='red')
```

120

Having clause: generalized OR

- `bool_or(test)`: true, if test is true for at least one line of group
- Supplier offering both green and red articles

```
select c.fid
from catalogue c, articles a
where a.aid=c.aid
group by c.fid
    having bool_or(acoul='vert')
        and
        bool_or(acoul='red')
```

121

Tuple construction

- The comparison with the embedded query can involve more than one attribute.
- The attributes must be enclosed by a pair of parentheses (tuple constructor)
- Our previous query can be rewritten as:

```
select *
from Person P
where (Name,LastName) in
    (select Name, LastName
     from Person P1
     where P1.NumSecu <> P.NumSecu)
```

122

Comments on subqueries

- Embedded queries can be ‘less declarative’, but are mostly easier to read
- Complex queries with variables can be hard to understand.
- The embedded queries can not contain set operations, mostly (take home lesson: “only do unions on top level”). This limitation is not significant, and not present in all DBMS.

123

Comments on subqueries

- The use of variables must respect rules of visibility
 - a variable can only be used in the query where it is introduced, or within subqueries embedded therein
 - If a variable name is ambiguous, the system assumes we are referring to the closer one

124

Visibility of variables

- Incorrect query:

```
select *
from Customer
where Cus_ID in
      (select Cus_ID
       from Contract O1
       where Con_ID = 'AZ1020')
or Cus_ID in
      (select Cus_ID
       from Contract O2
       where O2.Date = O1.Date)
```
- The query is incorrect, because the variable O1 is not visible within the second embedded query.

125

Subqueries in modification commands

Modification commands with `in`

- Increase by 5 euro the VALUE of all contracts that contain the product 456

```
update Contract
set VALUE = VALUE + 5
where Con_ID in
      (select Con_ID
       from Detail
       where Prod_ID = '456')
```

126

127

Embedded queries in modifications

- Setting TotalPieces to the sum of quantities of all lines of a contract.

```
update Contract O
  set TotalPieces =
    (select sum(Qt)
     from Detail D
     where D.Con_ID = O.Con_ID)
```

128

Views

- Offer the "view" of virtual tables (external schemas)
- Classified into:
 - simple (selection and projection from only one table)
 - complex

- Syntax:

```
create view ViewName [ (AttributeList) ]
  as Subquery
  [ with [ local | cascaded ] check option ]
```

130

Next topics

- Views
- Generic constraints
- Access control

129

Views

- Their definition may contain other views, that were previously defined, but without mutual dependency (recursion was introduced in SQL:1999)
- Can be used to write complex queries
 - Query decomposition
- Are sometimes needed to express certain queries
 - Namely such queries that combine and embed several aggregate operations

131

Composition of views and queries

- View creation:

```
create view MainContracts as
  select *
  from Contract
  where VALUE > 10000
```

- Query:

```
select Cus_ID
from MainContracts
```

- Composition of both:

```
select Cus_ID
from MainContracts
where VALUE > 100000
```

132

Views and queries

- Extract the customer with the highest total bill (without view):

```
select Cus_ID
from Contract
group by Cus_ID
having sum(VALUE) >= all
      (select sum(VALUE)
       from Contract
       group by Cus_ID)
```

- Works with Postgresql, but not accepted by all SQL systems.

133

Views and queries

- Extract the customer with the highest bill (via view):

```
create view CustomerBill(Cus_ID,TotalBill)
as
  select Cus_ID, sum(VALUE)
  from Contract
  group by Cus_ID;

select Cus_ID
from CustomerBill
where TotalBill = (select max(TotalBill)
                  from CustomerBill);
```

134

Views and queries

- Extract the average number of contracts per customer:

- Incorrect query (aggregate functions can not be nested):

```
select avg(count(*))
from Contract
group by Cus_ID
```

- Correct query (with a view):

```
create view CustomerStat=(Cus_ID,ConNumber) as
select Cus_ID, count(*)
from Contract
group by Cus_ID;
```

```
select avg(ConNumber)
from CustomerStat;
```

135

Example of simple view

- Contracts with VALUE over 10.000

```
create view MajorContracts as
  select *
  from Contract
  where VALUE > 10000
```

Contract

1	3	1-6-96	50.000
4	1	1-7-12	12.000
6	3	3-9-12	27.000

VIEW:
Major contracts

136

Simple views in a cascade

```
create view Administrators
  (Sid,Name,LastName,Income) as
select Sid, Name, LastName, Income
from Employee
where Department = 'Administration'
```

```
create view JuniorAdministrators as
select *
from Administrators
where Income < 50
with check option
```

137

Modifications through views

- View:

```
create view MajorContracts as
  select *
  from Contract
  where VALUE > 10000
```

- Modification:

```
update MajorContracts
  set VALUE = VALUE * 1.05
  where Cus_ID = '45'
```

- Composition of both:

```
update Contract
  set VALUE = VALUE * 1.05
  where Cus_ID = '45'
  and VALUE > 10000
```

138

Check option: updating views

- The **check option** acts when the content of a view is modified.
 - Pre-condition: inserted/ updated tuple must be part of the view.
 - Post-condition: the tuple must remain in the view
- local**: control only with respect to the view that is invoking the command.
- cascaded**: the control is made in all involved views, recursively.

139

Check option: example

- `create view MajorContracts70 as
select *
from MajorContracts
where Cus_ID = '70'
with local check option`
- **Dependencies:**
 - MajorContracts: Contracts with VALUE>10000
 - MajorContracts70: MajorContracts with Cus_ID=70

140

Check option

- `update MajorContracts70
set Cus_ID = '71'
where Con_ID = '754'`

is refused with check option **local** and **cascaded**
- `update MajorContracts70
set VALUE = 5000
where Con_ID = '754'`

is accepted with **local**, but refused with **cascaded**

141

Complex view


What else is possible, beyond selection and projection?

```
create view CusPro(Customer,Product) as  
select Cus_ID, Prod_ID  
from Contract join Detail  
on Contract.Con_ID = Detail.Con_ID
```

142

Complex view (JOIN)

Customer	Product
12	45



143

Query on complex view

- Query:


```
select Customer
from CusPro
where Product = '45'
```
- Combining both:


```
select Cus_ID
from Contract join Detail
on Contract.Con_ID = Detail.Con_ID
where Prod_ID = '45'
```

144

Modifications of the complex view

- It is impossible to modify the original table through the view, because the interpretation is ambiguous:
- Ex.:

```
update CusPro
set Product = '42'
where Customer = '12'
```
- Ambiguity for the modification of the original tables
 - The customer has changed his contract
 - The product's identifier has changed

145

Complex view (JOIN)

Customer	Product
12	45

42

JOIN

Cus_ID	Con_ID	Con_ID	Prod_ID
12	33	45	33	45	42
			45	42	

146

Recursion in SQL:1999

```
with recursive Raggiungibile (Orig, Dest, Costo) as
( select Orig, Dest, Costo
  from Volo where Orig = 'Milano'
  union
  select V.Orig, R.Dest, V.Costo+R.Costo
  from Volo V join Raggiungibile R
    on V.Dest = R.Orig
  where V.Dest not in
    select V.Dest in Raggiungibile)

select distinct Dest, Costo
from Raggiungibile R
where Costo = (select min(Costo)
               from Raggiungibile R1
               where R.Dest = R1.Dest)
```

147

Constraints in the Data Definition Language (DDL)

Generic integrity constraints

- Predicates that must hold on correct (legal) instances of the database
- Expressed in two ways:
 - in the table's schema
 - as separate assertions

Data quality

- Data quality
 - Correctness, completeness, up-to-date?
 - Quality of real data is often poor (5- 40% incorrect)
- To improve the data quality:
 - Integrity rules
 - Data manipulation by predefined programs (procedures and triggers)

Check clause

- Allows to express arbitrary constraints in the schema definition.
- It appears immediately after the attribute, within the **create table** command.
- Syntax:
 - check** (*Condition*)
- *Condition* is what can appear in a `where` clause (including embedded queries), i.e.its evaluation returns a boolean value

Example

- Employee(Emp_ID,FirstName,LastName,Dept,Superior)
 - Managers, whose ID starts with digit 1, may not have a superior
 - Otherwise, an employee's superior must be from the same department
- Example: constraints for the attribute *Superior* in the schema of the table *Employee*:

```
Superior character(6)
check (Emp_ID like "1%"
      or
      Department = (select Department
                     from Employee E
                     where E.Sid = Superior))
```

152

Assertions

- Assertions allow to define constraints outside of table definitions, by giving a name to a check clause
- Useful in many situations, for example, to express generic constraints between tables
- Syntax:

```
create assertion AssertionName check (Condition)
```
- Ex: the table *Employee* must contain at least one tuple:

```
create assertion AlwaysOneEmployee
check (1 <= (select count(*)
             from Employee))
```

153


When are constraints checked?

immediate :

violation cancels the last modification

deferred (later):

violation cancels the whole application



Dynamic modification of the meaning of constraints

- Each constraint is defined as of a certain type** (usually "immediate")
- The application can modify the initial type of constraints:
 - `- set constraints immediate`
 - `- set constraints deferred`
- Sooner or later, all constraints are checked.

155

Example: managing a shop

Shop

Prod_ID	QtDisp	QtOrder
1	150	100
3	130	80
4	170	50
5	500	150

Order

Prod_ID	Date	QtaOrd

156

Example: definition of the shop

```
create table Shop as
( Prod_ID    char(2) primary key,
  QtDisp     integer not null
              check (QtDisp > 0) ,
  QtOrd      integer not null
              check (QtaOrd > 10)
              check (QtaDisp>QtaOrd)
)
```

157

Access control

Access control

- Privacy: protection of the DB in order to guarantee that only authorized users may access it
- Mechanisms to identify the user (by *password*):
 - When she connects to the computer system
 - When she connects to the DBMS
- Individual users, and user groups

159

Permissions

- Each component of a scheme can be protected (tables, attributes, views, domains, etc)
- A resource's owner (its creator) assigns **privileges (permissions)** to other users
- A pre-defined user **_system** represents the administrator, and has full access to all resources
- A privilege is specified by:
 - The resource
 - The user giving the privilege
 - The user receiving the privilege
 - The action that is allowed on the resource
 - The possibility to pass on the permission to other users

160

Grant and revoke

- Syntax to give a privilege to a user:
`grant < Privileges | all privileges > on Resource
to User [with grant option]`
 - **grant option** indicates if the grant can be propagated to other users.
- To withdraw a privilege:
`revoke Privileges on Resource from User
[restrict | cascade]`

162

6 types of privileges in SQL

- **insert**: add a new object to the resource
- **update**: modify the resource's content
- **delete**: remove an object from the resource
- **select**: access the resource's content in queries
- **references**: create a referential integrity constraint that involves the resource (may restrict the possibility to modify the resource!)
- **usage**: use the resource in a schema definition (particularly, a domain)
- **all privileges**: summarizes all 6 types

161

Examples

```
grant all privileges on Contract to User1
grant update(VALUE) on Contract to User2
grant select on Contract to User2, User3
```

```
revoke update on Contract from User1
revoke select on Contract from User3
```

163

Example of grant option

1 Database administrator

```
grant all privileges on Contract to User1  
with grant option
```

2 User1

```
grant select on Contract to User2  
with grant option
```

3 User2

```
grant select on Contract to User3
```

164

Withdrawing a privilege with cascade

1 Database administrator

```
grant select on Contract to User1  
with grant option
```

2 User1


```
grant select on Contract to User2
```

3 Database administrator

```
revoke select on Contract from User1 cascade
```

165

Withdrawing a privilege with cascade



166

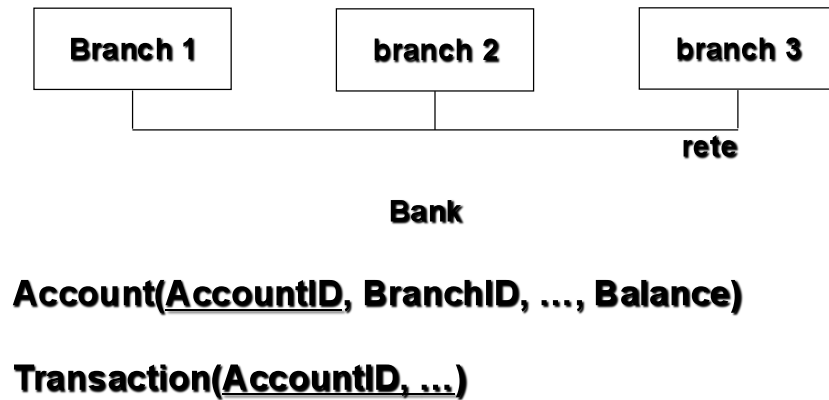
Views and access control

Views = unit of permission

- Allows the optimal management of privacy.

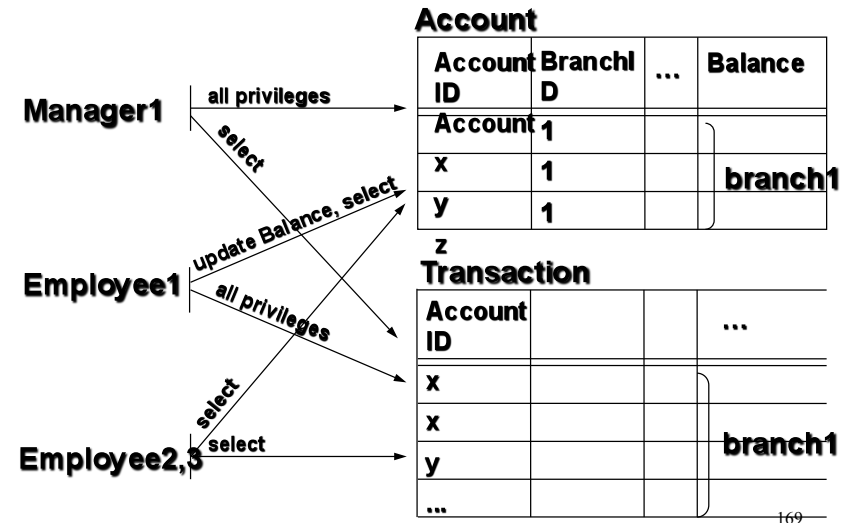
167

Example: managing bank accounts



168

Access needs



169

Views relative to the first branch

```
create view Account1 as
( select *
  from Account
  where branch = 1)

create view Transaction1 as
( select *
  from Transaction
  where AccountID in
    ( select AccountID
      from Account1 ) )
```

170

Permissions relative to data of the first branch

```
grant all privileges on Account1
  to Manager1
grant update(Balance) on Account1
  to Employee1
grant select on Account1
  to Employee1, Employee2, Employee3
grant select on Transaction1
  to Manager1
grant all privileges on Transaction1
  to Employee1
grant select on Transaction1
  to Employee2, Employee3
```

171